

Changed files

1. defs.h (line 123 [last line])

```
104 //PAGEBREAK: 16
105 // proc.c
106 int      cpuid(void);
107 void     exit(void);
108 int      fork(void);
109 int      growproc(int);
110 int      kill(int);
111 struct cpu* mycpu(void);
112 struct proc* myproc();
113 void     pinit(void);
114 void     procdump(void);
115 void     scheduler(void) __attribute__((noreturn));
116 void     sched(void);
117 void     setproc(struct proc*);
118 void     sleep(void*, struct spinlock*);
119 void     userinit(void);
120 int      wait(void);
121 void     wakeup(void*);
122 void     yield(void);
123 int      countTraps(void);
```

(line 123)

2. syscall.h

```
23 #define SYS_countTraps 22
```

3. user.h (line 26 [last line])

```
4 // system calls
5 int fork(void);
6 int exit(void) __attribute__((noreturn));
7 int wait(void);
8 int pipe(int*);
9 int write(int, const void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int);
13 int exec(char*, char**);
14 int open(const char*, int);
15 int mknod(const char*, short, short);
16 int unlink(const char*);
17 int fstat(int fd, struct stat*);
18 int link(const char*, const char*);
19 int mkdir(const char*);
20 int chdir(const char*);
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
26 int countTraps(void);
```

4. proc.h (lines 52-53 [last lines])

```
52 int countSyscall[22]; // Syscall counter
53 int countTraps[20]; // Trap counter
```

5. usys.S

```
32 SYSCALL(countTraps)
```

6. syscall.c

```
106 extern int sys_countTraps(void);
130 [SYS_countTraps] sys_countTraps,
```

7. sysproc.c

```
93 int
94 sys_countTraps(void)
95 {
96     return countTraps();
97 }
```

8. proc.c

```
9 #include "syscall.h" // include header of syscall
10 #include "traps.h" // include header of traps.h
```

```
12 struct {
13     struct spinlock lock;
14     struct proc proc[NPROC];
15 } ptable;
16
17 /**
18 * Initializes both countSyscall and countTraps arrays
19 * Assigns each element of the arrays to be 0
20 */
21 void assignArr(struct proc *p) {
22     for(int i = 0; i < 22; i++) {
23         p->countSyscall[i] = 0;
24     }
25     for(int j = 0; j < 20; j++) {
26         p->countTraps[j] = 0;
27     }
28 }
```

(lines 21-82)

```
29
30 /**
31 * array of names of syscall commands
32 */
33 static char *syscallNameCode[] = {
34     [SYS_fork] "SYS_fork",
35     [SYS_exit] "SYS_exit",
36     [SYS_wait] "SYS_wait",
37     [SYS_pipe] "SYS_pipe",
38     [SYS_read] "SYS_read",
39     [SYS_kill] "SYS_kill",
40     [SYS_exec] "SYS_exec",
41     [SYS_fstat] "SYS_fstat",
42     [SYS_chdir] "SYS_chdir",
43     [SYS_dup] "SYS_dup",
44     [SYS_getpid] "SYS_getpid",
45     [SYS_sbrk] "SYS_sbrk",
46     [SYS_sleep] "SYS_sleep",
47     [SYS_uptime] "SYS_uptime",
48     [SYS_open] "SYS_open",
49     [SYS_write] "SYS_write",
50     [SYS_mknod] "SYS_mknod",
51     [SYS_unlink] "SYS_unlink",
52     [SYS_link] "SYS_link",
53     [SYS_mkdir] "SYS_mkdir",
54     [SYS_close] "SYS_close",
55     [SYS_countTraps] "SYS_countTraps"
56 };
```

```
57
58 /**
59 * array of names of traps
60 */
61 static char *trapNameCode[] = {
62     [T_DIVIDE] "T_DIVIDE",
63     [T_DEBUG] "T_DEBUG",
64     [T_NMI] "T_NMI",
65     [T_BRKPT] "T_BRKPT",
66     [T_OFLOW] "T_OFLOW",
67     [T_BOUND] "T_BOUND",
68     [T_ILLOP] "T_ILLOP",
69     [T_DEVICE] "T_DEVICE",
70     [T_DBLFLT] "T_DBLFLT",
71     [T_TSS-1] "T_TSS",
72     [T_SEGNP-1] "T_SEGNP",
73     [T_STACK-1] "T_STACK",
74     [T_GPFLT-1] "T_GPFLT",
75     [T_PGFLT-1] "T_PGFLT",
76     [T_FPERR-2] "T_FPERR",
77     [T_ALIGN-2] "T_ALIGN",
78     [T_MCHK-2] "T_MCHK",
79     [T_SIMDERR-2] "T_SIMDERR",
80     [18] "T_SYSCALL",
81     [19] "T_DEFAULT"
82 };
```

```
179 sp = sizeof *p->context;
180 p->context = (struct context*)sp;
181 memset(p->context, 0, sizeof *p->context);
182
183 // assigns array of p
184 assignArr(p);
185
186 p->context->eip = (uint)forkret;
187
188 return p;
189 }
190
191 //PAGEBREAK: 32
192 // Set up first user process.
193 void
194 userinit(void)
```

(Line 184; at the bottom of method 'alloproc();)

```

325 // Parent might be sleeping in wait().
326 wakeup1(curproc->parent);
327
328 /**
329 * adds up both syscalls and trap calls
330 * from parent and child process
331 */
332 int i;
333 for(i = 0; i < 22; i++) {
334     curproc->parent->countSyscall[i] = curproc->parent->countSyscall[i] + curproc->countSyscall[i];
335     if(i < 20) {
336         curproc->parent->countTraps[i] = curproc->parent->countTraps[i] + curproc->countTraps[i];
337     }
338 }
339
340 // Pass abandoned children to init.
341 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
342     if(p->parent == curproc){
343         p->parent = initproc;
344         if(p->state == ZOMBIE)
345             wakeup1(initproc);
346     }
347 }

```

(lines 332-338; bottom of 'exit()' method)

```

621 /**
622 * This function is our system call. It keeps track of both system calls and traps calls.
623 * At first, it counts their occurrence from the current process.
624 * -> This is saved in countTraps[] and countSyscall[] arrays.
625 * Then, it displays them to the shell screen.
626 * After that, it shows the names of each call in both countTraps[] and countSyscall[] arrays,
627 * and the number of their occurrences in a process next to them.
628 * At last, it shows how many times the user process has been trapped to the OS in general.
629 */
630 int
631 countTraps(void) {
632     int totalSyscall = 0;
633     int totalTraps = 0;
634
635     for(int i = 0; i < 20; i++) {
636         totalTraps += myproc()->countTraps[i];
637     }
638
639     printf("\n*-----*\n");
640     printf("| Total number of traps = %d\t|\n", totalTraps);
641     printf("|-----|\n");
642     printf("| Trap Name \t| Occurrences\t|\n");
643     printf("|-----|\n");
644
645     for(int j = 0; j < 20; j++) {
646         if(myproc()->countTraps[j] != 0) {
647             printf("| [%s]\t| %d\t|\n", trapNameCode[j], myproc()->countTraps[j]);
648         }
649     }
650     printf("*-----*\n");
651
652     for(int m = 0; m < 22; m++) {
653         totalSyscall += myproc()->countSyscall[m];
654     }
655
656     printf("\n*-----*\n");
657     printf("| Total number of system calls = %d\t|\n", totalSyscall);
658     printf("|-----|\n");
659     printf("| Syscall Name \t\t| Occurrences\t|\n");
660     printf("|-----|\n");
661

```

```

662 for(int n = 0; n < 22; n++) {
663     if(myproc()->countSyscall[n] != 0) {
664         if(n == 21) { // if the call is [SYS_countTraps] ; does a bit of formatting
665             cprintf("| [%s] \t| %d\t\t|\n", syscallNameCode[n + 1], myproc()->countSyscall[n]);
666         }
667         else {
668             cprintf("| [%s] \t\t| %d\t\t|\n", syscallNameCode[n + 1], myproc()->countSyscall[n]);
669         }
670     }
671 }
672 cprintf("*-----*\n");
673
674 cprintf("\n*-----*\n");
675 if(totalTraps + totalSyscall < 10) {
676     cprintf("| Total number of all calls = %d\t\t|\n", totalTraps + totalSyscall);
677 }
678 else {
679     cprintf("| Total number of all calls = %d\t|\n", totalTraps + totalSyscall);
680 }
681 cprintf("*-----*\n\n");
682
683 return 22;
684 }

```

9. Trap.c

```

11 // Interrupt descriptor table (shared by all CPUs).
12 struct gatedesc idt[256];
13 extern uint vectors[]; // in vectors.S: array of 256 entry pointers
14 struct spinlock tickslock;
15 uint ticks;
16
17 /**
18 * Checks which trap is called based on the current process's trap number.
19 * The function then increments that trap number by one each time to give us information at the end about
20 * what types of traps have occurred and the number of times they (traps) have occurred.
21 */
22 void
23 check_trap_func(struct trapframe *tf)
24 {
25     if(tf->trapno == T_DIVIDE) {
26         myproc()->countTraps[0] = myproc()->countTraps[0] + 1;
27     }
28     else if(tf->trapno == T_DEBUG) {
29         myproc()->countTraps[1] = myproc()->countTraps[1] + 1;
30     }
31     else if(tf->trapno == T_NMI) {
32         myproc()->countTraps[2] = myproc()->countTraps[2] + 1;
33     }
34     else if(tf->trapno == T_BRKPT) {
35         myproc()->countTraps[3] = myproc()->countTraps[3] + 1;
36     }
37     else if(tf->trapno == T_OFLOW) {
38         myproc()->countTraps[4] = myproc()->countTraps[4] + 1;
39     }
40     else if(tf->trapno == T_BOUND) {
41         myproc()->countTraps[5] = myproc()->countTraps[5] + 1;
42     }
43     else if(tf->trapno == T_ILLOP) {
44         myproc()->countTraps[6] = myproc()->countTraps[6] + 1;
45     }
46     else if(tf->trapno == T_DEVICE) {
47         myproc()->countTraps[7] = myproc()->countTraps[7] + 1;
48     }
49     else if(tf->trapno == T_DBLFLT) {
50         myproc()->countTraps[8] = myproc()->countTraps[8] + 1;
51     }

```

```

52  else if(tf->trapno == T_TSS) {
53      myproc()->countTraps[9] = myproc()->countTraps[9] + 1;
54  }
55  else if(tf->trapno == T_SEGNP) {
56      myproc()->countTraps[10] = myproc()->countTraps[10] + 1;
57  }
58  else if(tf->trapno == T_STACK) {
59      myproc()->countTraps[11] = myproc()->countTraps[11] + 1;
60  }
61  else if(tf->trapno == T_GPFLT) {
62      myproc()->countTraps[12] = myproc()->countTraps[12] + 1;
63  }
64  else if(tf->trapno == T_PGFLT) {
65      myproc()->countTraps[13] = myproc()->countTraps[13] + 1;
66  }
67  else if(tf->trapno == T_FPEERR) {
68      myproc()->countTraps[14] = myproc()->countTraps[14] + 1;
69  }
70  else if(tf->trapno == T_ALIGN) {
71      myproc()->countTraps[15] = myproc()->countTraps[15] + 1;
72  }
73  else if(tf->trapno == T_MCHK) {
74      myproc()->countTraps[16] = myproc()->countTraps[16] + 1;
75  }
76  else if(tf->trapno == T_SIMDERR) {
77      myproc()->countTraps[17] = myproc()->countTraps[17] + 1;
78  }
79  else if(tf->trapno == T_SYSCALL) {
80      myproc()->countTraps[18] = myproc()->countTraps[18] + 1;
81  }
82  else if(tf->trapno == T_DEFAULT) {
83      myproc()->countTraps[19] = myproc()->countTraps[19] + 1;
84  }
85 }
86
87 void
88 tvinit(void)
89 {

```

(lines 22-85)

```

105 //PAGEBREAK: 41
106 void
107 trap(struct trapframe *tf)
108 {
109     if(tf->trapno == T_SYSCALL){
110         if(myproc()->killed)
111             exit();
112
113         /** Both system call and trap number are handled.
114          * If a trap number exists, then check_trap_func()
115          * is invoked and the new trap call increments
116          * the occurrence of that trap by 1 in array countTraps[20].
117          * If the eax is valid, then the system call occurrence
118          * is incremented by 1 in array countSyscall[22].
119          */
120         int num;
121         int i;
122         num = tf->eax;
123         i = num - 1;
124         myproc()->tf = tf;
125         myproc()->countSyscall[i] = myproc()->countSyscall[i] + 1;
126         check_trap_func(myproc()->tf);
127
128         syscall();
129         if(myproc()->killed)
130             exit();
131         return;
132     }

```

(line 120 - 126)

10. Makefile (lines 184-185 ; line 255)

```

166 .PRECIOUS: %.o
167
168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184     _test_case1\
185     _test_case2\
186
187 fs.img: mkfs README $(UPROGS)
188     ./mkfs fs.img README $(UPROGS)
189
190 -include *.d
191
192 clean:

```

```

246 # CUT HERE
247 # prepare dist for students
248 # after running make dist, probably want to
249 # rename it to rev0 or rev1 or so on and then
250 # check in that version.
251
252 EXTRA=\
253     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
254     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\
255     printf.c umalloc.c test_case1.c test_case2.c\
256     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
257     .gdbinit.tmpl gdbutil\
258
259 dist:
260     rm -rf dist
261     mkdir dist
262     for i in $(FILES); \
263     do \
264         grep -v PAGEBREAK $$i >dist/$$i; \
265     done
266     sed '/CUT HERE, $$d' Makefile >dist/Makefile
267     echo >dist/runoff.spec
268     cp $(EXTRA) dist

```