

Timur Gaimakov: A20415319

Hassan Alamri: A204730047

## Files modified

*syscall.h* (lines 23-24)

```
22 #define SYS_close 21
23 #define SYS_GetSharedPage 22
24 #define SYS_FreeSharedPage 23
25
```

*usys.S* (lines 32-33)

```
31 SYSCALL(uptime)
32 SYSCALL(GetSharedPage)
33 SYSCALL(FreeSharedPage)
```

*user.h* (lines 26-27)

```
25 int uptime(void);
26 void* GetSharedPage(int, int);
27 int FreeSharedPage(int);
28
29 // ulib.c
30 int stat(const char*, struct stat*);
```

*proc.h* (lines 38-41 and line 58)

```
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
37 //Added: tracks mapped shared pages
38 struct shared_mapping {
39     int key;
40     void *virt_addr;
41 };
42
43 // Per-process state
44 struct proc {
45     char name[16]; // Process name (debugging)
46     struct shared_mapping shared_memory[64]; // Process shared pages
47 };
48
49 // Process memory is laid out contiguously, low addresses first:
```

*syscall.c* (lines 106-107 and 131-132)

```
105 extern int sys_uptime(void);
106 extern int sys_GetSharedPage(void);
107 extern int sys_FreeSharedPage(void);
108
109 static int (*syscalls[])(void) = {
110     [SYS_fork] sys_fork,
111     [SYS_close] sys_close,
112     [SYS_GetSharedPage] sys_GetSharedPage,
113     [SYS_FreeSharedPage] sys_FreeSharedPage,
114 };
115 void
```

*proc.c* (lines 116-119; line 183 [between growproc() and fork()]; line 209 [in fork()])

```
113 p->context->eip = (uint)forkret;
114
115 // Added: initializes values of shared mapping
116 for(int i = 0; i < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); i++) {
117     p->shared_memory[i].key = -1;
118     p->shared_memory[i].virt_addr = 0;
119 }
120 return p;
121 }
122
123 //PAGEBREAK: 32
124
125 }
126
127 //Added: copies shared mapping from one region to another
128 extern int copy_shared_memory_regions(struct proc *, struct proc *);
129
130 // Create a new process copying p as the parent.
131 // Sets up stack to return as if from system call.
132 // Caller must set state of returned proc to RUNNABLE.
133 int
134 fork(void)
```

```

207
208 //Added: copies shared pages from curproc to np
209 copy_shared_memory_regions(curproc, np);
210
211 np->sz = curproc->sz;

```

**sysproc.c** (EOF; lines 93 - 120)

```

90     return xticks;
91 }
92
93 extern void *GetSharedPage(int i, int len);
94
95 void*
96 sys_GetSharedPage(void)
97 {
98     int key;
99     int num_of_pages;
100
101     if(argint(0, &key) < 0){
102         return (void*)-1;
103     }
104     else if(argint(1, &num_of_pages) < 0){
105         return (void*)-1;
106     }
107     return (void*)(GetSharedPage(key, num_of_pages));
108 }
109
110 extern int FreeSharedPage(int id);
111
112 int
113 sys_FreeSharedPage(void)
114 {
115     int key;
116
117     if(argint(0, &key) < 0){
118         return -1;
119     }
120     return FreeSharedPage(key);
121 }

```

**vm.c** (lines 13, 285, 302, 403 - EOF)

```

10 extern char data[]; // defined by kernel.ld
11 pde_t *kpgdir; // for use in scheduler()
12
13 int is_PA_linked_with_sharedMemory(uint);
14
284 // Added
285 void forced_free_phys_page(void *);
286
287 // Free a page table and all the physical memory pages
288 // in the user part.
289 void
290 freevm(pde_t *pgdir)

```

```

300
301 //Added
302 forced_free_phys_page(v);
303 }
304 }
305 kfree((char*)pgdir);
306 }

```

```

401 // Blank page.
402
403 // Added
404 #define MAX_REGION_SIZE 64
405 typedef char bool;
406 struct shared_memory_region {
407     bool valid;
408     int len;
409     int ref_count;
410     uint phys_pages[MAX_REGION_SIZE];
411 };
412
413 struct shared_memory_region regions[64];
414
415 /**
416 * Maps region "key" into process "p"
417 * at virtual address "addr"
418 */
419 void map_shared_memory_region(int key, struct proc *p, void *virt_addr) {
420     for(int k = 0; k < regions[key].len; k++) {
421         mappages(p->pgdir, (void*)(virt_addr + (k * PGSIZE)), PGSIZE, regions[key].phys_pages[k], PTE_M | PTE_U);
422     }
423 }

```

```

425  /** Syscall that returns a virt_addr
426  * to the shared page when successful and the process
427  * can start writing and reading from the virt_addr.
428  */
429  void *
430  GetSharedPage(int key, int num_of_pages)
431  {
432      /**
433       * Check if key is valid
434       */
435      if(key < 0 || key > 64){
436          return (void*)0;
437      }
438      /**
439       * If key is invalid, then allocate virtual pages
440       * into appropriate physical pages' region
441       */
442      else if(regions[key].valid == 0) {
443          for(int j = 0; j < num_of_pages; j++) {
444              void* new_page = kalloc(); // Allocate new page
445              memset(new_page, 0, PGSIZE); // Empty page
446              regions[key].phys_pages[j] = V2P(new_page); // Save that new physical page
447          }
448          regions[key].valid = 1;
449          regions[key].len = num_of_pages;
450          regions[key].ref_count = 0;
451      }
452      /**
453       * If regions' length does not equal
454       * to number of pages, then exit
455       */
456      else if(regions[key].len != num_of_pages){
457          return (void*)0;
458      }
459      regions[key].ref_count++;
460
461      // Get process's "p" index
462      struct proc *p = myproc();
463      int index = -1;
464      for(int i = 0; i < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); i++) {
465          if(p->shared_memory[i].key == -1) {
466              index = i;
467              break;
468          }
469      }
470      /**
471       * Pages is filled with 0's initially
472       */
473      if(index == -1){
474          return (void*)0;
475      }
476
477      // Get the least VAS (virt_addr_space) currently used
478      void *virt_addr = (void*)(KERNBASE - PGSIZE);
479      for(int j = 0; j < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); j++) {
480          if(p->shared_memory[j].key != -1 &&
481             ( (uint)(virt_addr) > (uint)(p->shared_memory[j].virt_addr) ) ) {
482              virt_addr = p->shared_memory[j].virt_addr;
483          }
484      }
485
486      // Get VA (virt_addr) of new assigned pages
487      virt_addr = (void*)virt_addr - (PGSIZE * num_of_pages);
488      p->shared_memory[index].virt_addr = virt_addr;
489      p->shared_memory[index].key = key;
490
491      // Map that virt_addr in memory
492      map_shared_memory_region(key, p, virt_addr);
493
494      return virt_addr;
495  }

```

```

497 /** Copy the shared memory regions
498  * of process "p" into process "new_p"
499  */
500 int copy_shared_memory_regions(struct proc *p, struct proc *new_p) {
501     for(int i = 0; i < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); i++) {
502         if(p->shared_memory[i].key != -1) {
503             new_p->shared_memory[i] = p->shared_memory[i];
504             int key = new_p->shared_memory[i].key;
505             regions[key].ref_count++;
506             map_shared_memory_region(key, new_p, new_p->shared_memory[i].virt_addr);
507         }
508     }
509     return 0;
510 }
511
512 /** Syscall that removes the calling process
513  * from accessing the shared pages
514  * associated with the key.
515  */
516 int
517 FreeSharedPage(int key)
518 {
519     // Free shared memory structure
520     struct proc *p = myproc();
521     void *virt_addr = 0;
522     for(int i = 0; i < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); i++){
523         if(p->shared_memory[i].key == key) {
524             virt_addr = p->shared_memory[i].virt_addr;
525             p->shared_memory[i].key = -1;
526             p->shared_memory[i].virt_addr = 0;
527             break;
528         }
529     }
530     if(virt_addr == 0){
531         return -1;
532     }
533
534     // Clear page table entries (pte)
535     struct shared_memory_region* reg = &regions[key];
536
537     for(int i = 0; i < reg->len; i++) {
538         pte_t* pte = walkpgdir(p->pgdir, (char*)virt_addr + i * PGSIZE, 0);
539         if(pte == 0) {
540             return -1;
541         }
542         *pte = 0;
543     }
544
545     // Decrease the reference count; free if ref_count isn't used
546     reg->ref_count--;
547     if(reg->ref_count == 0) {
548         regions[key].valid = 0;
549         regions[key].ref_count = 0;
550         for(int i = 0; i < regions[key].len; i++){
551             kfree(P2V(regions[key].phys_pages[i]));
552         }
553         regions[key].len = 0;
554     }
555     return 0;
556 }

```

```

558 /**
559  * Check if given PA (phys_addr)
560  * is linked with a shared memory page
561  */
562 int is_PA_linked_with_sharedMemory(uint phys_addr){
563     struct proc *p = myproc();
564     for(int i = 0; i < (int)(sizeof(p->shared_memory) / sizeof(p->shared_memory[0])); i++) {
565         for(int j = 0; j < regions[i].len; j++) {
566             if(phys_addr == (uint)regions[i].phys_pages[j]) {
567                 return 1;
568             }
569         }
570     }
571     return 0;
572 }
573
574 /**
575  * Free pages when
576  * is_PA_linked_with_sharedMemory (are not shared)
577  * fails (returns 0)
578  */
579 void forced_free_phys_page(void *phys_addr) {
580     if(is_PA_linked_with_sharedMemory(V2P(phys_addr)) == 0){
581         kfree(phys_addr);
582     }
583 }

```

*tester.c* (create file)

```
1 #include "types.h"
2 #include "user.h"
3
4 int main() {
5     void* region = GetSharedPage(5, 3);
6     for(int i = 0; i < 10; i++) {
7         printf(1, "%d ", ((char*)region)[i]);
8     }
9     printf(1, "\n");
10
11     // write
12     strcpy(region, "region");
13
14     // read
15     printf(1, "%s\n", region);
16
17     // FreeSharedPage(0);
18     exit();
19 }
```

*Makefile*

```
166 .PRECIOUS: %.o
167 # Added: tester
168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184     _tester\
```